

Replication Certification at Publication— Checkable by Anyone, Anywhere, in Seconds, Forever*

Gary King[†] Yiqing Xu[‡] Leo Y. Yang[§]

October 1, 2026

Abstract

Authors increasingly deposit replication datasets—the data and code behind published results—in archives such as Dataverse, and journals increasingly require it. But a deposit alone proves little, so authors keep paying: months while a data editor reruns analyses, years answering readers’ questions, and rancor as failed replications become accusations because no record establishes what actually ran. We introduce a *certificate*: a cryptographically signed record of the deposited files and output captured by specialized logging software. It replaces the data editor’s rerun with evidence of the author’s own run that anyone can check, which no one, including the archive, can backdate, and no one can fake without risking detection. Certification asks little of anyone and benefits everyone: The author deposits, then reruns the analysis one extra time; the archive automatically hashes, signs, timestamps it in Bitcoin’s blockchain, and stores the record—confidential data included, under restricted access, where licenses allow—but never runs the analysis; the journal simply checks a website that the certificate exists; and anyone can verify the record in seconds, or replay any step of it, even decades later.

*We thank Katalina Toth and Salil Vadhan for helpful comments and suggestions. The present version of this paper and the Supplementary Appendix are available at GaryKing.org/certify.

[†]Institute for Quantitative Social Science, Harvard University; GaryKing.org, King@Harvard.edu.

[‡]Department of Political Science, Stanford University; yiqingxu.org, yiqingxu@stanford.edu.

[§]Stanford Center on China’s Economy and Institutions, Stanford University, leoyang@stanford.edu; currently affiliated with the Department of Accountancy, Economics and Finance, Hong Kong Baptist University; LeoYang.org, leoyang@hkbu.edu.hk.

Contents (Page to be removed before publication)

1	Introduction	3
2	Problems with Replication Policies	3
2.1	Policy Enforcement Levels	4
2.2	Policy Failures	4
2.3	Policy Costs	5
3	A Technological Solution: The Certificate	6
3.1	How It Works	6
3.2	Why A Certificate Is Evidence, Not Merely a Promise	9
3.3	What Everyone Gains	9
4	What a Certificate Does Not Establish	11
5	Prior Work	11
6	Conclusion	12

1. Introduction

Ensuring that evidence offered in a scholarly publication stands independently of its author is a goal that echoes through the history of science, back even to the 1660 motto of the Royal Society, *Nullius in verba* (“take nobody’s word for it”). We can also see this goal in the now widely accepted “*replication standard* [which] holds that sufficient information exists with which to understand, evaluate, and build upon a prior work if a third party could replicate the results without any additional information from the author,” (King, 1995, p.444), and in many efforts to meet the standard.¹

For this purpose, scholarly publication outlets over four decades have adopted increasingly strict policies designed to convince authors to deposit *replication datasets*, the data and code behind a published result, in professional archives and to try to ensure each deposit provides meaningful evidence that results can be reproduced. Unfortunately, each of the policy changes has merely moved the burden of verifying replicability from one actor to another, while often increasing costs in money, staffing, publication delays, time investment by authors and future researchers, and misunderstandings between them, while still never quite accomplishing the goal (see Section 2).

Instead of even stricter and more expensive policies, we propose a technological solution aimed at accomplishing the same goals more effectively at lower costs. By marshalling well-tested cryptographic methods, we show how to create a special replication dataset *certificate* that provides the evidence that the analysis ran on exactly its files and produced exactly the results it records, with its random draws made after the files were fixed. The journal can then substitute its expensive and time-consuming in-house replication runs with a quick check that the certificate exists before final acceptance, and a later reader can check the record in seconds instead of re-running the analysis. A failed replication then has a certified record to be argued from rather than a reputation to be argued about (see Section 3).

The rest of the paper states what a certificate does not establish (Section 4) and what our proposal borrows from prior work and what that work leaves open (Section 5). The Supplementary Appendix provides all technical details, including what evading certification costs an attacker (Section S5.2) and our contributions (Section S6).

2. Problems with Replication Policies

No journal policy tried thus far lets anyone verify a deposited replication dataset without re-running it, so future researchers pay repeatedly to verify the same deposit. We describe here the levels of enforcement journals have escalated through; the evidence indicating how these policies have failed; and the substantial costs imposed by the policies in time, money, and rancor.²

¹We define “replication” as most users and archives do: running the same code on the same data, producing the same result. Some call this “reproduction” and reserve “replication” for a study redone with new data (National Academies of Sciences, Engineering, and Medicine, 2019); we use the words interchangeably.

²The norm that authors deposit the code and data behind a published result is long-established in the social sciences (Dewald, Thursby and Anderson, 1986; King, 1995), and professional archives now supply the storage, versioning, and persistent identifiers this protocol builds on (King, 2007). The Transparency and Openness Promotion guidelines extend the same levels across the sciences, with verified re-execution as their

2.1. Policy Enforcement Levels

Level 1 is the original and least stringent proposal, where journals require authors to “indicate in their first footnote in which public archive they will deposit the information necessary to replicate their numerical results,” and journals would police that footnote and nothing more. Editors would be “responsible only for the existence of the footnote, not for confirming that the data set has been submitted nor for checking whether the results actually can be replicated.” Any verification “can and should be left to future researchers” (King, 1995).

Level 2 requires authors to go a step further and make data available on request, such as *Science*’s rule from 2011. But asked for the materials behind 204 papers published under it, authors supplied them for only 44%, and the results of only 26% of the 204 could be reproduced (Stodden, Seiler and Ma, 2018).

Level 3 makes deposit of the data and code a condition of publication, as now required by many leading outlets (Brodeur et al., 2026; Ferguson et al., 2023; Miske et al., 2026). In economics, the requirement reached the *American Economic Review* in 2005 and the *Quarterly Journal of Economics* in 2016 (Christensen and Miguel, 2018). In political science, the editors of twenty-seven journals signed the joint statement of the Data Access and Research Transparency initiative, committing to require that cited data be available in a trusted repository (Lupia and Elman, 2014; Data Access and Research Transparency (DART), 2015). Among quantitative political-science articles, a working link to the data and code was rare before about 2008, with the share growing approximately exponentially since (Rainey et al., 2025, fig.1).

Finally, Level 4 is re-execution of the code by the journal, originally as an option and then a requirement. From July 2009, *Biostatistics* let authors ask it to run their code and marked the article *R* if it reproduced; in two years, five of 125 articles earned the mark (Peng, 2011). Journals now increasingly require human-verified re-execution by the journal’s data editor (Vilhuber, 2020; Nüst and Eglen, 2021). The *American Journal of Political Science* worked with the Odum Institute, a university data archive and verification service, to verify every accepted quantitative article since 2015, with a mean verification time near seventy-five days (Jacoby, 2018).

2.2. Policy Failures

The literature suggests three conclusions. First, *archival deposit alone does not show that the code runs*. Of more than nine thousand R scripts deposited in Dataverse, 74% did not run to completion without error (Trisovic et al., 2022). Of 15,817 Python notebooks from biomedical publications with declared dependencies, 1,203 of those ran through without error (Samuel and Mietchen, 2024), in line with earlier mass executions of notebooks (Pimentel et al., 2019). When two teams each tried to redo one analysis from each of 18 microarray papers, ten could not be reproduced (Ioannidis et al., 2009).

Second, *running the code does not guarantee that the materials reproduce the article’s results*. Working only from deposited files, one audit reproduced the key result of just 22 of 67 published economics papers (Chang and Li, 2022). At *Cognition*, mandated deposit raised the availability of usable data sharply, while 13 of 35 audited articles still had at least one

top level (Nosek et al., 2015).

value that could not be reproduced even with the authors’ help (Hardwicke et al., 2018). In the largest journal-run audit, 68% of attempts succeeded in replicating results (Fišar et al., 2024).

And finally, *journal verification catches more failures without catching them all*. Of the twenty-four papers given in-house replication review at the *Quarterly Journal of Political Science*, 14 of 18 produced results that differed from those in the manuscript (Eubank, 2016). Across 254 papers from the three leading political-science journals whose materials could be obtained, 61 did not reproduce every coefficient exactly and 20 did not reproduce it even approximately (Xu and Yang, 2026).

2.3. Policy Costs

Finally, the policies impose three categories of cost: time, money, and rancor. The cost in time comes in two parts, delayed acceptance and endless post-publication questions. In the first, median verification times at the American Economic Association’s journals have run three months per round, and some manuscripts need several rounds (Vilhuber, 2024); the average manuscript the *AER* published in 2025 spent thirteen weeks in the data editor’s hands, up from nine in 2022 (Luttmer, 2026). The second time cost is that authors assemble replication datasets to a data editor’s standard but still have to answer questions about them for years. For later researchers, re-using a deposited replication dataset is far better than no deposit but still considerable work: in one mass-reproduction project, teams averaged thirteen days per paper, and about one in twelve spent more than a hundred days (Brodeur et al., 2024). None of that work is recorded where the next reader can use it and so every researcher starts from scratch.

The second cost is paid by journal editors who hire and manage replicators, in addition to providing cutting edge computational resources used by researchers. The American Economic Association’s Data Editor pays for all student employment and pre-doctoral staff, while Cornell University’s Economics Department provides tuition waivers and computing (Vilhuber, 2025). The Cornell laboratory that replicates received, in one year, 529 requests covering 415 manuscripts. About 15-20 undergraduate research assistants work there per week, supervised by the three staff (Vilhuber et al., 2022). At the *American Journal of Political Science*, verification by the Odum Institute took an average of eight person-hours per manuscript, and authors resubmitted materials 1.7 times on average, a “not trivial” financial cost (Jacoby, Lafferty-Hess and Christian, 2017). *Management Science* charges \$79 on every original submission since 2025, to fund its reproducibility infrastructure and external verification (INFORMS, 2025).

The final cost is rancor. When a replication fails, the absence of a shared record of what was actually run often turns technical questions into personal ones. The dispute over Reinhart and Rogoff’s result turned on what had actually been run—which spreadsheet, which exclusions, which weighting—and was settled only when the authors found and shared the original file (Herndon, Ash and Pollin, 2014). The dispute over LaCour and Green’s article turned on whether the deposited data had been collected at all (Broockman, Kalla and Aronow, 2015). Daniel Kahneman, acting as a senior statesman, has gone so far as to propose etiquette rules for how readers attempting replications should approach authors (Kahneman, 2014).

3. A Technological Solution: The Certificate

Figure 1 gives an overview of certification, with this section structured corresponding to each of its three panels: how certification works (Section 3.1); why it is sufficient evidence (Section 3.2); and what everyone gains (Section 3.3).

3.1. How It Works

Certification has three parts corresponding to the three bands within Panel 1 of Figure 1. First, the author makes the replication dataset **READY** by locking down all sources of variability (due to the code, computing environment, or settings), so that it gives the same results if run repeatedly. Second, the author and archive work together to **CERTIFY** the replication dataset. And finally, anyone, at any point in the future, can **CHECK** the results in seconds without having to rerun the analysis.

Ready: author prepares the replication dataset. The author uses our certificate software to ensure their replication dataset gives the same results when run twice on their machine, hence defining “replicability” for this study. The software first collects (a) every file in the replication dataset, with its hash, and the results of a *dependency audit*, a reading of the scripts for the files they open and the seeds they set; (b) the settings to freeze, such as the time zone and the number of processor cores (Section S2.2); and (c) the software environment itself (deposited as a “container image”, a packaged copy of the operating system and programs, where software licenses allow, and otherwise as add-on libraries installed inside the replication dataset; Section S2.1). From these, the software creates a *closed copy*, which includes the replication dataset plus the small files added by the certificate software. The closed copy enables all the author’s analyses to be run with one top-level command, keeps track of each step in the process, and, where possible, has only one input for a random seed.

Finally, the author conducts two runs of the closed copy, what we call the *closure runs*, each from a fresh temporary copy and under a seed of the author’s choosing. The certificate software compares every file the two runs wrote and grades their agreement on a scale of six *determinism tiers*, from byte-for-byte identical, through equal within a stated numeric tolerance, to a real difference or a run that did not finish (Table S1). The dataset is *ready* when the two runs agree to within half a unit in the fourth significant figure of every number, though the goal is byte-identical agreement; the software writes a report, deposited with the replication dataset, that names any file that differed (Section S2.2).

Certify: author and archive fix the files, run once, seal the record. In three steps (corresponding to the Certify row of Panel 1 in Figure 1), we show how to take a dataset that is “ready” and certify it (Section S3 gives each step in full). The steps are designed to eliminate the possibility of *backdating* by anyone including even the archive (in order to surreptitiously change the record after the fact), *seed shopping* (where the author cherry-picks a supposedly random seed to make their results look better), and *ambiguity* (in what results the run produces or how to conduct a new one).

First, the author *deposits* the closed copy in the archive. The archive then computes a new random seed with three inputs: a hash of the uploaded files, the archive’s secret




Figure 1. Certification at a glance. Panel 1: **READY**, the author prepares the replication files until two runs agree. **CERTIFY**, the archive issues a seed for the fingerprinted files, the author runs the analysis under it, and the archive signs and timestamps the record; draws fall after the seed is issued and before the certificate is issued. **CHECK**, three options, depending on need. Panel 2: why the record is evidence. Panel 3: what each party pays **TODAY** and gains once **CERTIFIED**.

cryptographic key, and a recent Bitcoin block — meaning that no one, the author included, could know the seed before it is computed. The archive’s receipt, its signed and timestamped statement that these files arrived, records the hash, the block, and the seed (Section S3.1).

Second, the author *runs* the closed copy once under the seed issued by the archive (producing the “certified run”). For each step in the run, the certificate software records the “pre-state” (every file and its hash present before the step starts), any files it writes, and any output it produces. The published numbers must come from this run (Section S3.2). The software produces a “proposed certificate”.

Finally, the author submits the proposed certificate for the archive to *seal*. The archive checks the certificate, refuses on any failure, signs the “witness” (its statement that it sealed this certificate, using its cryptographic key), timestamps the certificate on the blockchain, and enters the witness in its ledger (a public list of every document it issues including this one, kept so that any later edit is visible publicly). The archive stores the certificate and records it in a metadata field linked to the published dataset (Section S3.3; Section S8 lists everything an archive must do).

Check: anyone can easily verify the record. Anyone can check a certificate on their own computer, without taking anyone else’s word for what was run or when (Section S4). Three types of checks are available, depending on need.

First, the *existence check* enables a journal to confirm that the replication dataset, the article’s citation of it, and the certificate exist, and that the certificate carries no exception mark (the archive’s flag indicating that the certificate falls short of one of the protocol’s standards). This requires only a fast lookup on the archive’s webpage; no reanalysis is necessary (Section S4.3). An existence check is sufficient for the journal because a sealed certificate enables anyone, without help from the journal, to verify both “record consistency” (that the stored files are those the run started from and wrote, and that neither the files nor the record has changed since) and “replication” (that rerunning the recorded steps from their stored starting files produces outputs that agree with the certified ones within the certificate’s determinism tier).

Second, the *record check* is designed for someone who prefers not to take the archive’s webpage on trust. Here, our open-source certificate software recomputes every hash from the files and checks the archive’s signatures, the timestamps, and the ledger, validating record consistency. The record check runs no analysis, typically takes from seconds to a few minutes, needs no outside service when run offline, and, as long as the cryptography it uses stays infeasible to break, works the same way decades from now (Sections S4.1 and S5.1.4).

Finally, the *re-execution check* is designed for a future researcher who wishes to build on this research or to test whether the record was faked. The researcher can rerun the entire analysis and see whether its results agree with the certified ones within the determinism tier. A rerun takes as long as the original, is one command where a container image was deposited, and is judged only when the researcher’s software environment matches the recorded one (Section S4.2); a rerun that agrees gives the researcher a tested starting point for new work. Conveniently, a researcher can also rerun a single step, without running any of those that came before or after, whenever all of that step’s starting files are stored; the certificate reports how many steps qualify. For confidential data, only readers the archive approves can run the

parts that read the files, but the certificate is still issued and published (Section S3.5).

3.2. Why A Certificate Is Evidence, Not Merely a Promise

Unlike articles published under prior replication policies, a certified paper carries its own evidence, which anyone can check, anywhere, in seconds, forever, without taking anyone’s word for what was run or when: not the author’s, not the journal’s, not even the archive’s. *Evidence* here means a record anyone can check, which no one, including the archive, can backdate, and no one can fake without risking detection. This is summarized in four features (corresponding to Panel 2 of Figure 1).

First, *replication dataset files are fixed before the certification run*. The archive hashes the deposit and computes the seed from that hash, so the seeded draws cannot precede the files; this rests on well-tested cryptography, not on trust (Section S1). Seed shopping therefore needs a new deposit for each seed, and the archive’s ledger publicly records every deposit (Section S5.1).

Second, *the results are fixed after the certification run*. The certificate records the hash of every file each step the code started from and produced and every logged estimate, and the witness names the certificate’s own hash. Thus, a later change to any file or to the record fails the check.

Third, *the certification run is fixed in time*, meaning that a researcher trying to correct earlier mistakes will fail or leave a visible trail. The receipt and the certificate are each timestamped in Bitcoin’s blockchain, so the certified run’s random draws are bounded between two block times. The later edge is the time of the block that timestamps the certificate. The earlier edge is the time of the block the receipt names, since the seed is computed from that block’s hash (Section S3.3.2). Not even the archive can backdate either timestamp, and that earliest date is what no signature, log, or timestamp could supply on its own (Section 5).

Finally, *the result is checkable by anyone*. No special permission is required (except to obtain access if the replication dataset is confidential). Everything else the checks test is arithmetic, which anyone can easily recompute starting from the replication dataset and the certificate. The archive’s word is taken only for two items: that it issued the seed as specified (Section S7) and the facts only it holds (who deposited and whether a revised deposit and an earlier deposit it names as its predecessor have the same uploader; see Section S3.4). That the run happened as recorded rests on the open-source certificate software, which any reader can test by re-running recorded steps (Section S4.2).

3.3. What Everyone Gains

Certification lowers the cost of verification by changing who runs the analysis and how often: the author runs it once more, on the machine that already holds their software environment, and everyone afterward checks the record instead of rebuilding and rerunning it from scratch. Panel 3 of Figure 1 shows the gains party by party. Because several gains reach more than one party, we take them here by the problem in Section 2 that each one solves, with what certification costs in the same terms. Those problems are the failures of Section 2.2 and the three costs of Section 2.3, taken in that order – policy failures and costs in time, money, and rancor.

First, policy failures. Current replication datasets, even when required by journal policies, carry no evidence that their code runs or reproduces the article’s results, and journal verification catches only some issues (Section 2.2). In contrast, a certified replication dataset carries evidence that its analysis ran on exactly its files and produced exactly the results it records, which the copyeditor or author checks during copy editing to connect the certificate to the publication (Section S4.3). Because checking time does not grow with the length of the computation (Section S4.1), the evidence extends to analyses too long for a data editor to re-execute, which happen because published articles often represent the computational cutting edge. The closure runs, of the kind authors already make before submission anyway, are now deposited and fixed in the record, where readers can test them, while today no one else ever sees an author’s own runs.

Second, time. Data editors presently must rebuild the author’s software environment, rerun their analyses, and ask the author for whatever is missing. This takes months per round before acceptance, each later researcher rebuilds it all from scratch, and authors answer questions about their deposits for years (Section 2.3). Certification replaces publication delays and later questions with one more run by the author at final acceptance, on the machine that already holds the hardware, software, libraries, and data, so that the run comes down to nearly one command. A future reader or researcher attempting a replication confirms that the record is intact with the record check, which is fast and requires no copy of the author’s software environment (Section 3.1), and answers from the record the questions about what was run that would otherwise go to the author. Confirming that a step still produces its recorded outputs takes that step’s own running time (where the container image was deposited; without an image, the reader must also install the recorded software versions; Section S4.2). The author pays in making one extra run; a revision that changes the files after certification repeats all three runs (Section S3.4). The author of a restricted deposit also checks before release that nothing the certificate publishes identifies a person or organization in those files (Section S3.5).

Third, money. Journals today pay for the data editors, staff, and computing that re-execution requires (Section 2.3). For a certified paper, journal staff only need an existence check (webpage lookup) at final acceptance with no technical choices (or special expertise) required, because the protocol fixes the standard every certificate must meet. Archives pay to operate the software that performs the deposit and seal steps, and for storage of the record, which includes the files the run writes and so can exceed the deposit itself (Section S3.6). The archive’s institutional work of custody, versions, backups, metadata, terms of use, restricted access, and indefinite storage continues, as before.

Finally, rancor. Today a failed replication becomes a dispute about what the author did, because no shared record says what was run (Section 2.3). Certification gives the scientific community a public record of precisely what was run for every certified paper, so a dispute over a failed replication turns on what the record shows, which both sides can check, instead of on what the author says was run. The publication and its certified replication dataset then stand independently of the person who created them. Disputes will continue, but only about future research, as should be the case.

4. What a Certificate Does Not Establish

Certification provides evidence only of record consistency and replication — that both the inputs and the outputs are what is claimed. No check catches code that fabricates its results, although its files and date are fixed in a public record that cannot be backdated. Nor does any check guarantee that the dependency audit finds every file the code reads or every seed it sets (Section S5.1). Record checks do not prevent an author from surreptitiously altering a run through administrator or hardware control of the computer to change results, although such behavior can be caught by a re-execution check. And even re-execution checks cannot guarantee that the results answer the paper’s questions, avoids p-hacking, reflect patterns in freshly collected data, or is otherwise good science — as this is where the replication standard finishes and novel advances in science begin.

Certification requires that the archive hold the data: data licensed under terms that forbid deposit with a third party, or usable only inside a secure facility, cannot be certified, and papers using them still need re-execution by someone with access to the data (Section S3.5). Procedurally, we rely on the copyeditor or author to tie the printed page to outputs of the certified run. Our protocol trusts the archive that it issued the seed as specified and for the account facts only it holds; a deposit shows who held the bytes, not who wrote them. A stolen signing key yields certificates with valid signatures that are missing from the ledger, which the check exposes once the reader holds a later summary of the ledger that the archive cannot revise; a leaked copy of the key that computes seeds would make future seeds predictable without a trace; and a compromised signing program that also controls the ledger can pass every check, but cannot backdate. A modified build of the certificate software that reports the expected hash would not be caught.

Record consistency can be checked indefinitely, but replication decays as the recorded software goes out of service. A reader can then emulate the recorded environment, or rewrite the analysis and compare results with the certified estimates, except where a proprietary package can no longer be licensed or installed (Section S5.1.4).

The supplementary appendix gives every limit in full (Section S5, especially Section S5.1). We provide a framework for studying what intentional attempts to evade certification cost in the supplementary appendix (Section S5.2).

5. Prior Work

We build on three bodies of prior work (Section S6 summarizes our contributions): (1) designs that give evidence about a published analysis; (2) the tools the protocol borrows; and (3) methods that date work by drawing public randomness after a commitment.

First, the designs, in order of increasing assurance. File fingerprints show whether a file changed ([The Dataverse Project, 2026](#)), and verifiable computational results store each published result under an identifier the article cites ([Gavish and Donoho, 2011](#)). Signed provenance, graded by the framework SLSA (Supply-chain Levels for Software Artifacts), signed execution records that authors can neither alter nor disown, and platforms such as Whole Tale that package the environment for re-execution rest on the word of whoever signed or packaged ([Open Source Security Foundation, 2026](#); [Keita and Homan, 2026](#); [Brinckman et al., 2019](#)). TRACE, whose goal is closest to ours, requires its own platform to run and

attest to the workflow, which also lets it reach data usable only inside a secure facility (Vilhuber et al., 2023; Li et al., 2025). Journal data editors, the named checkers of CODECHECK, and the French certification agency cascadi re-run the analysis themselves, at the cost of a qualified person per paper (Nüst and Eglen, 2021; Pérignon et al., 2019). Trusted hardware signs output produced inside a protected processor, on the word of the chip’s maker (Costan and Devadas, 2016; Li, 2026); cryptographic proofs of computation need no such trust but remain far too costly today for general data analysis (Goldwasser, Micali and Rackoff, 1989; Walfish and Blumberg, 2015). A certificate needs neither a platform that vouches for the run nor a second party’s re-run.

Second, the borrowed tools. The cryptography is well-tested: Merkle trees, keyed hashes, digital signatures, and public timestamps (Section S7), obtained in Bitcoin through OpenTimestamps (Todd, 2016), with a ledger modeled on the append-only logs that Certificate Transparency keeps for website certificates and Sigstore for signed software, an internet standard (Laurie, Langley and Kasper, 2013; Newman, Meyers and Torres-Arias, 2022; Birkholz et al., 2026). Blockchain timestamps of research files and pipelines exist, but show only that files existed by some time (Gipp, Meuschke and Gernandt, 2015; Wittek et al., 2021; Casas and Fidler, 2024). The recording of a run draws on running twice under frozen settings and comparing the outputs (Lamb and Zacchiroli, 2022; Bizopoulos and Bizopoulos, 2020), and on signed per-step records of inputs and outputs and the provenance workflow engines keep (Torres-Arias et al., 2019). Re-executing a random sample of recorded steps comes from distributed computing (Monrose, Wyckoff and Rubin, 1999; Du et al., 2004; Haeberlen et al., 2010; Jia et al., 2021); the deterrence random checks provide has been analyzed as well (Golle and Mironov, 2001; Barnett, Zardo and Graves, 2018).

Finally, dating work by public randomness drawn after a “commitment”, a public statement that fixes the materials. Election audits (Cordero, Wagner and Dill, 2006; Lindeman and Stark, 2012), the Internet Engineering Task Force’s selection of committee members (Eastlake, 2004), randomness beacons (Rabin, 1983; Kelsey et al., 2019), and a randomization seeded from a Bitcoin block (Saa and Stern, 2019) all draw their randomness after publishing such a commitment, and a signature can postdate a blockchain’s latest block, which dates the signature rather than a run’s draws (Abadi et al., 2020). Seeds hashed from the commitment itself (Fiat and Shamir, 1987) or from a preregistration (Mazor, Mazor and Mukamel, 2019), and blind analysis, which hides results under a perturbation of the data until the analysis is frozen (MacCoun and Perlmutter, 2015), leave trust to whoever commits it, a gap that running the whole study once inside protected hardware closes, at the price of ruling out re-execution (Meißner et al., 2021). A certificate needs no separately published commitment.

6. Conclusion

A certificate is designed to move the scientific community closer to the idea of keeping evidence in publications based solely on the record and independent of the author. It turns one additional run of the analysis by the author already prepared to do so, made under the archive’s seed after deposit, into evidence that anyone can check, and it lets a journal replace its own rerun with a simple look-up on a webpage. We suggest the journal print the replication dataset’s usual citation, by DOI and version, with one word added, “Certified”

(Section S4.3). The supplementary appendix gives the more specific specification for version 1 of our protocol, from how an author makes a dataset ready (Section S2) to the byte-level rules an implementer needs (Section S9).

References

- Abadi, Aydin, Michele Ciampi, Aggelos Kiayias and Vassilis Zikas. 2020. Timed Signatures and Zero-Knowledge Proofs—Timestamping in the Blockchain Era. In *Applied Cryptography and Network Security (ACNS 2020)*. Lecture Notes in Computer Science Springer pp. 335–354.
- Barnett, Adrian G., Pauline Zardo and Nicholas Graves. 2018. “Randomly Auditing Research Labs Could Be an Affordable Way to Improve Research Quality: A Simulation Study.” *PLOS ONE* 13(4):e0195613.
- Birkholz, Henk, Antoine Delignat-Lavaud, Cédric Fournet, Yogesh Deshpande and Steve Lasker. 2026. “An Architecture for Trustworthy and Transparent Digital Supply Chains.” RFC 9943, RFC Editor. <https://www.rfc-editor.org/rfc/rfc9943>
- Bizopoulos, Paschalis and Dimitris Bizopoulos. 2020. “Reconciler: A Workflow for Certifying Computational Research Reproducibility.” arXiv:2005.12660v3.
- Brinckman, Adam, Kyle Chard, Niall Gaffney, Mihael Hategan, Matthew B. Jones, Kacper Kowalik, Sivakumar Kulasekaran, Bertram Ludäscher, Bryce D. Mecum, Jarek Nabrzyski, Victoria Stodden, Ian J. Taylor, Matthew J. Turk and Kandace Turner. 2019. “Computing Environments for Reproducibility: Capturing the “Whole Tale”.” *Future Generation Computer Systems* 94:854–867. <https://doi.org/10.1016/j.future.2017.12.029>
- Brodeur, Abel, Derek Mikola, Nikolai Cook et al. 2024. Mass Reproducibility and Replicability: A New Hope. Discussion Paper 16912 IZA Institute of Labor Economics. <https://docs.iza.org/dp16912.pdf>
- Brodeur, Abel, Derek Mikola, Nikolai Cook et al. 2026. “Reproducibility and Robustness of Economics and Political Science Research.” *Nature* 652:151–156. <https://doi.org/10.1038/s41586-026-10251-x>
- Broockman, David, Joshua Kalla and Peter Aronow. 2015. “Irregularities in LaCour (2014).” Stanford University. https://stanford.edu/~dbroock/broockman_kalla_aronow_lg_irregularities.pdf
- Casas, Santiago and Christian Fidler. 2024. “TOPO: Time-Ordered Provable Outputs.” arXiv:2411.00072 [astro-ph.IM]. <https://arxiv.org/pdf/2411.00072>
- Chang, Andrew C. and Phillip Li. 2022. “Is Economics Research Replicable? Sixty Published Papers from Thirteen Journals Say ‘Often Not’.” *Critical Finance Review* 11(1):185–206. <https://doi.org/10.1561/104.00000053>

- Christensen, Garret and Edward Miguel. 2018. “Transparency, Reproducibility, and the Credibility of Economics Research.” *Journal of Economic Literature* 56(3):920–980. <https://doi.org/10.1257/jel.20171350>
- Cordero, Arel, David Wagner and David Dill. 2006. The Role of Dice in Election Audits—Extended Abstract. In *IAVoSS Workshop on Trustworthy Elections (WOTE 2006)*. <https://people.eecs.berkeley.edu/~daw/papers/dice-wote06.pdf>
- Costan, Victor and Srinivas Devadas. 2016. “Intel SGX Explained.” IACR Cryptology ePrint Archive, Report 2016/086. <https://eprint.iacr.org/2016/086>
- Data Access and Research Transparency (DA-RT). 2015. “Data Access and Research Transparency (DA-RT): A Joint Statement by Political Science Journal Editors.” *State Politics & Policy Quarterly* 15(4):423–424. <https://doi.org/10.1177/1532440015598318>
- Dewald, William G., Jerry G. Thursby and Richard G. Anderson. 1986. “Replication in Empirical Economics: The Journal of Money, Credit and Banking Project.” *American Economic Review* 76(4):587–603. <https://www.jstor.org/stable/1806061>
- Du, Wenliang, Jing Jia, Manish Mangal and Mummoorthy Murugesan. 2004. Uncheatable Grid Computing. In *Proceedings of the 24th International Conference on Distributed Computing Systems (ICDCS 2004)*. pp. 4–11.
- Eastlake, 3rd, Donald E. 2004. “Publicly Verifiable Nominations Committee (NomCom) Random Selection.” RFC 3797, Internet Engineering Task Force.
- Eubank, Nicholas. 2016. “Lessons from a Decade of Replications at the Quarterly Journal of Political Science.” *PS: Political Science & Politics* 49(2):273–276. https://www.nickeubank.com/wp-content/uploads/2019/10/Eubank_LessonsFromADecadeOfReplications.pdf
- Ferguson, Joel, Rebecca Littman, Garret Christensen, Elizabeth Levy Paluck, Nicholas Swanson, Zenan Wang, Edward Miguel, David Birke and John-Henry Pezzuto. 2023. “Survey of Open Science Practices and Attitudes in the Social Sciences.” *Nature Communications* 14:5401. <https://doi.org/10.1038/s41467-023-41111-1>
- Fiat, Amos and Adi Shamir. 1987. How to Prove Yourself: Practical Solutions to Identification and Signature Problems. In *Advances in Cryptology—CRYPTO ’86*. Vol. 263 of *Lecture Notes in Computer Science* Springer pp. 186–194.
- Fišar, Miloš, Ben Greiner, Christoph Huber, Elena Katok, Ali I. Ožkeš and the Management Science Reproducibility Collaboration. 2024. “Reproducibility in Management Science.” *Management Science* 70(3):1343–1356. <https://doi.org/10.1287/mnsc.2023.03556>
- Gavish, Matan and David L. Donoho. 2011. “A Universal Identifier for Computational Results.” *Procedia Computer Science* 4:637–647.
- Gipp, Bela, Norman Meuschke and André Gernandt. 2015. Decentralized Trusted Timestamping using the Crypto Currency Bitcoin. In *Proceedings of the iConference 2015*. Newport Beach, CA: . <https://arxiv.org/pdf/1502.04015>

- Goldwasser, Shafi, Silvio Micali and Charles Rackoff. 1989. “The Knowledge Complexity of Interactive Proof Systems.” *SIAM Journal on Computing* 18(1):186–208. <https://doi.org/10.1137/0218012>
- Golle, Philippe and Ilya Mironov. 2001. Uncheatable Distributed Computations. In *Topics in Cryptology — CT-RSA 2001*. Springer pp. 425–440. https://doi.org/10.1007/3-540-45353-9_31
- Haeblerlen, Andreas, Paarijaat Aditya, Rodrigo Rodrigues and Peter Druschel. 2010. Accountable Virtual Machines. In *Proceedings of the 9th USENIX Symposium on Operating Systems Design and Implementation (OSDI '10)*. pp. 119–134.
- Hardwicke, Tom E., Maya B. Mathur, Kyle MacDonald, Gustav Nilsson, George C. Banks, Mallory C. Kidwell, Alicia Hofelich Mohr, Elizabeth Clayton, Erica J. Yoon, Michael Henry Tessler, Richie L. Lenne, Sara Altman, Bria Long and Michael C. Frank. 2018. “Data Availability, Reusability, and Analytic Reproducibility: Evaluating the Impact of a Mandatory Open Data Policy at the Journal Cognition.” *Royal Society Open Science* 5(8):180448. <https://doi.org/10.1098/rsos.180448>
- Herndon, Thomas, Michael Ash and Robert Pollin. 2014. “Does High Public Debt Consistently Stifle Economic Growth? A Critique of Reinhart and Rogoff.” *Cambridge Journal of Economics* 38(2):257–279. <https://doi.org/10.1093/cje/bet075>
- INFORMS. 2025. “Submission Fee FAQ.” Management Science. <https://pubsonline.informs.org/page/mnsc/submission-fee-faq>
- Ioannidis, John P. A., David B. Allison, Catherine A. Ball, Issa Coulibaly, Xiangqin Cui, Aedín C. Culhane, Mario Falchi, Cesare Furlanello, Laurence Game, Giuseppe Jurman, Jon Mangion, Tapan Mehta, Michael Nitzberg, Grier P. Page, Enrico Petretto and Vera van Noort. 2009. “Repeatability of published microarray gene expression analyses.” *Nature Genetics* 41(2):149–155. <https://doi.org/10.1038/ng.295>
- Jacoby, William G. 2018. “The American Journal of Political Science Replication and Verification Policy.” American Journal of Political Science. <https://ajps.org/wp-content/uploads/2019/03/ajps-replic-and-verif-policy-2-27-18.pdf>
- Jacoby, William G., Sophia Lafferty-Hess and Thu-Mai Christian. 2017. “Should Journals Be Responsible for Reproducibility?” Inside Higher Ed. <https://www.insidehighered.com/blogs/rethinking-research/should-journals-be-responsible-reproducibility>
- Jia, Hengrui, Mohammad Yaghini, Christopher A. Choquette-Choo, Natalie Dullerud, Anvith Thudi, Varun Chandrasekaran and Nicolas Papernot. 2021. Proof-of-Learning: Definitions and Practice. In *2021 IEEE Symposium on Security and Privacy (S&P)*. pp. 1039–1056. <https://arxiv.org/abs/2103.05633>
- Kahneman, Daniel. 2014. “A New Etiquette for Replication.” *Social Psychology* 45(4):310–311. <https://doi.org/10.1027/1864-9335/a000202>

- Keita, Mamadou K. and Christopher Homan. 2026. “Computer Science Conferences Should Require Nonrepudiable Experimental Results.” arXiv:2605.08586.
- Kelsey, John, Luís T. A. N. Brandão, Rene Peralta and Harold Booth. 2019. A Reference for Randomness Beacons: Format and Protocol Version 2. Technical Report NIST IR 8213 (Initial Public Draft) National Institute of Standards and Technology. <https://nvlpubs.nist.gov/nistpubs/ir/2019/NIST.IR.8213-draft.pdf>
- King, Gary. 1995. “Replication, Replication.” *PS: Political Science & Politics* 28(3):444–452. <https://gking.harvard.edu/files/replication.pdf>
- King, Gary. 2007. “An Introduction to the Dataverse Network as an Infrastructure for Data Sharing.” *Sociological Methods & Research* 36(2):173–199. <https://gking.harvard.edu/files/dvn.pdf>
- Lamb, Chris and Stefano Zacchiroli. 2022. “Reproducible Builds: Increasing the Integrity of Software Supply Chains.” *IEEE Software* 39(2):62–70. <https://arxiv.org/pdf/2104.06020>
- Laurie, Ben, Adam Langley and Emilia Kasper. 2013. “Certificate Transparency.” RFC 6962, RFC Editor. <https://www.rfc-editor.org/rfc/rfc6962.txt>
- Li, Jiasun. 2026. “Trusted-Execution Environment (TEE) for Solving the Replication Crisis in Academia.” arXiv:2603.24878.
- Li, Meng, Timothy McPhillips, Craig Willis, Nikolaus Parulian, Bertram Ludäscher, Kacper Kowalik, Lars Vilhuber, Thu-Mai Lewis and Mandy Gooch. 2025. “TROV—A Model and Vocabulary for Describing Transparent Research Objects.” *International Journal of Digital Curation* 19(1):7.
- Lindeman, Mark and Philip B. Stark. 2012. “A Gentle Introduction to Risk-Limiting Audits.” *IEEE Security & Privacy* 10(5):42–49.
- Lupia, Arthur and Colin Elman. 2014. “Openness in Political Science: Data Access and Research Transparency.” *PS: Political Science & Politics* 47(1):19–42. <https://doi.org/10.1017/S1049096513001716>
- Luttmer, Erzo F. P. 2026. “Report of the Editor, American Economic Review.” *AEA Papers and Proceedings* 116:742–760. <https://doi.org/10.1257/pandp.116.742>
- MacCoun, Robert and Saul Perlmutter. 2015. “Blind Analysis: Hide Results to Seek the Truth.” *Nature* 526:187–189. <https://doi.org/10.1038/526187a>
- Mazor, Matan, Noam Mazor and Roy Mukamel. 2019. “A Novel Tool for Time-Locking Study Plans to Results.” *European Journal of Neuroscience* 49(9):1149–1156. <https://matanmazor.github.io/files/papers/mazor2019novel.pdf>
- Meißner, Echo, Felix Engelmann, Frank Kargl and Benjamin Erb. 2021. PeQES: A Platform for Privacy-Enhanced Quantitative Empirical Studies. In *Proceedings of the 36th ACM/SIGAPP Symposium on Applied Computing (SAC ’21)*. pp. 1226–1234. <https://arxiv.org/pdf/2103.05544>

- Miske, Olivia et al. 2026. “Investigating the reproducibility of the social and behavioural sciences.” *Nature* 652(8108):126–134. <https://doi.org/10.1038/s41586-026-10203-5>
- Monrose, Fabian, Peter Wyckoff and Aviel D. Rubin. 1999. Distributed Execution with Remote Audit. In *Proceedings of the Network and Distributed System Security Symposium (NDSS '99)*. <https://www.ndss-symposium.org/wp-content/uploads/2017/09/Distributed-Execution-with-Remote-Audit.pdf>
- National Academies of Sciences, Engineering, and Medicine. 2019. *Reproducibility and Replicability in Science*. Washington, DC: The National Academies Press. <https://doi.org/10.17226/25303>
- Newman, Zachary, John Speed Meyers and Santiago Torres-Arias. 2022. Sigstore: Software Signing for Everybody. In *Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security (CCS)*. pp. 2353–2367. <https://doi.org/10.1145/3548606.3560596>
- Nosek, Brian A., George Alter, George C. Banks et al. 2015. “Promoting an Open Research Culture.” *Science* 348(6242):1422–1425. <https://doi.org/10.1126/science.aab2374>
- Nüst, Daniel and Stephen J. Eglén. 2021. “CODECHECK: an Open Science initiative for the independent execution of computations underlying research articles during peer review to improve reproducibility.” *F1000Research* 10:253. Version 2. <https://f1000research.com/articles/10-253/v2/pdf>
- Open Source Security Foundation. 2026. “SLSA: Supply-chain Levels for Software Artifacts.” <https://slsa.dev/>.
- Peng, Roger D. 2011. “Reproducible Research in Computational Science.” *Science* 334(6060):1226–1227. <https://doi.org/10.1126/science.1213847>
- Pérignon, Christophe, Kamel Gadouche, Christophe Hurlin, Roxane Silberman and Eric Debonnel. 2019. “Certify reproducibility with confidential data.” *Science* 365(6449):127–128. <https://doi.org/10.1126/science.aaw2825>
- Pimentel, João Felipe, Leonardo Murta, Vanessa Braganholo and Juliana Freire. 2019. A Large-Scale Study About Quality and Reproducibility of Jupyter Notebooks. In *2019 IEEE/ACM 16th International Conference on Mining Software Repositories (MSR)*. pp. 507–517. <https://doi.org/10.1109/MSR.2019.00077>
- Rabin, Michael O. 1983. “Transaction Protection by Beacons.” *Journal of Computer and System Sciences* 27(2):256–267. [https://doi.org/10.1016/0022-0000\(83\)90042-9](https://doi.org/10.1016/0022-0000(83)90042-9)
- Rainey, Carlisle, Harley Roe, Qing Wang and Hao Zhou. 2025. “Data and Code Availability in Political Science Publications from 1995 to 2022.” *PS: Political Science & Politics* 58(2):339–345. <https://doi.org/10.1017/S1049096524001276>

- Saa, Olivia and Julio Michael Stern. 2019. Auditable Blockchain Randomization Tool. In *Proceedings of the 39th International Workshop on Bayesian Inference and Maximum Entropy Methods in Science and Engineering (MaxEnt 2019)*. Vol. 33 of *Proceedings MDPI* p. 17.
- Samuel, Sheeba and Daniel Mietchen. 2024. “Computational reproducibility of Jupyter notebooks from biomedical publications.” *GigaScience* 13:giad113. <https://doi.org/10.1093/gigascience/giad113>
- Stodden, Victoria, Jennifer Seiler and Zhaokun Ma. 2018. “An Empirical Analysis of Journal Policy Effectiveness for Computational Reproducibility.” *Proceedings of the National Academy of Sciences* 115(11):2584–2589. <https://doi.org/10.1073/pnas.1708290115>
- The Dataverse Project. 2026. “Universal Numerical Fingerprint (UNF).” <https://guides.dataverse.org/en/latest/developers/unf/index.html>.
- Todd, Peter. 2016. “OpenTimestamps: Scalable, Trust-Minimized, Distributed Timestamping with Bitcoin.” <https://petertodd.org/2016/opentimestamps-announcement>.
- Torres-Arias, Santiago, Hammad Afzali, Trishank Karthik Kuppusamy, Reza Curtmola and Justin Cappos. 2019. in-toto: Providing Farm-to-Table Guarantees for Bits and Bytes. In *28th USENIX Security Symposium*. USENIX Association pp. 1393–1410. <https://www.usenix.org/system/files/sec19-torres-arias.pdf>
- Trisovic, Ana, Matthew K. Lau, Thomas Pasquier and Mercè Crosas. 2022. “A Large-Scale Study on Research Code Quality and Execution.” *Scientific Data* 9:60. <https://www.nature.com/articles/s41597-022-01143-6.pdf>
- Vilhuber, Lars. 2020. “Reproducibility and Replicability in Economics.” *Harvard Data Science Review* 2(4). <https://hdsr.mitpress.mit.edu/pub/fgmpj11/download/pdf>
- Vilhuber, Lars. 2024. “Report of the AEA Data Editor.” *AEA Papers and Proceedings* 114:878–890. <https://doi.org/10.1257/pandp.114.878>
- Vilhuber, Lars. 2025. “Impacts of the AEA Verification Process.” Office of the AEA Data Editor. https://aeadataeditor.github.io/aea-cumulative-summary/impacts_of_aea_data_editing.html
- Vilhuber, Lars, Hyuk Harry Son, Meredith Welch, David N. Wasser and Michael Darisse. 2022. “Teaching for Large-Scale Reproducibility Verification.” *Journal of Statistics and Data Science Education* 30(3):274–281. <https://doi.org/10.1080/26939169.2022.2074582>
- Vilhuber, Lars, Thu-Mai Christian, Michał Kowalik, Bertram Ludäscher, Timothy McPhillips and Craig Willis. 2023. “TRACE: TRAnsparency CErtified — Trusting Computational Research Without Repeating It.” Zenodo. <https://doi.org/10.5281/zenodo.8000023>

- Walfish, Michael and Andrew J. Blumberg. 2015. “Verifying Computations Without Reexecuting Them.” *Communications of the ACM* 58(2):74–84. <https://dl.acm.org/doi/pdf/10.1145/2641562>
- Wittek, Kevin, Neslihan Wittek, James Lawton, Iryna Dohndorf, Alexander Weinert and Andrei Ionita. 2021. A Blockchain-Based Approach to Provenance and Reproducibility in Research Workflows. In *2021 IEEE International Conference on Blockchain and Cryptocurrency (ICBC)*. pp. 1–6.
- Xu, Yiqing and Leo Yang Yang. 2026. “Scaling Reproducibility: An AI-Assisted Workflow for Large-Scale Replication and Reanalysis.” arXiv:2602.16733. <https://arxiv.org/pdf/2602.16733>